

Computational Statistics - Homework 1

Leon Löppert, Jan Parlesak, Paul Jarschke

04.06.2024

Problem 1

Problem 1.1: Conjugate Gradient Algorithm

```
cg <- function(A, b, x) {  
  # Initializations:  
  r <- b - A %*% x # initial residual vector  
  p <- r # direction vector  
  j <- 0 # iteration counter  
  
  conv <- c() # convergence criterion tracker  
  err <- c() # error tracker  
  
  conv[1] <- norm(r, "2")  
  
  if (all(b == 0)) {  
    # if b is a zero vector, track error  
    err[1] <- norm(x, "I")  
  }  
  
  # Iterate until residual norm is sufficiently small  
  # or reached max number of iterations  
  cat("Starting Conjugate Gradient algorithm...\n")  
  while ((norm(r, "2") / norm(b, "2") > 1e-14) && j < 500) {  
    alpha <- (crossprod(r)) / (t(p) %*% A %*% p) # step size  
    x <- x + c(alpha) * p # Update solution vector x  
    rnew <- r - c(alpha) * (A %*% p) # new residual vector  
    beta <-  
      crossprod(rnew) / crossprod(r) # calculate beta coefficient  
    p <- rnew + c(beta) * p # update direction vector  
    r <- rnew # update residual vector  
    j <- j + 1 # update iteration counter  
    conv[j] <- norm(r, "2") / norm(b, "2") # track convergence  
  
    if (all(b == 0)) {  
      # if b is a zero vector, track error  
      err[j] <- norm(x, "I")  
    }  
  
    # Iteration counter:  
    cat(sprintf('It. %3.0f: %20.16e\n', j, conv[j]), "\n")  
  }
```

```

}
cat("...finished!\n")
cat("The solution for x is:", x, "\n")

# Convergence plots (Convergence criterion over iterations):
par(mfrow = c(2, 1))

plot(
  1:j,
  conv,
  type = "o",
  col = "blue",
  pch = 16,
  cex = 0.5,
  main = "Convergence Plot",
  xlab = "Iteration",
  ylab = "Convergence Criterion"
)
lines(1:j, conv, lty = 2, col = 'blue')

plot(
  1:j,
  conv,
  log = "y",
  type = "o",
  col = "red",
  pch = 16,
  cex = 0.5,
  main = "Convergence Plot",
  xlab = "Iteration",
  ylab = "Log Convergence Criterion"
)
lines(1:j, conv, lty = 2, col = 'red')

# Return the convergence history and the final solution vector
return(list(conv = conv, x = c(x)))
}

```

Problem 1.2: Preconditioned Conjugate Gradient Algorithm

```

pcg <- function(A, b, x, M, tol = 1e-14) {
  # Initializations: ----
  r <- b - A %*% x # initial residual vector
  z <- solve(M, r) # apply preconditioner M to the residual vector
  p <- z # direction vector
  j <- 1 # iteration counter

  conv <- c() # convergence criterion tracker
  err <- c() # error tracker

  conv[1] <- norm(r, "2")

```

```

if (all(b == 0)) {
  # if b is a zero vector, track error
  err[1] <- norm(x, "I")
}

# Iterate until residual norm is sufficiently small
# or reached max number of iterations
cat("Starting Preconditioned Conjugate Gradient algorithm...\n")

while ((norm(r, "2") / norm(b, "2") > tol) && j < 500) {
  alpha <- crossprod(z, r) / (t(p) %*% A %*% p) # step size
  x <- x + c(alpha) * p # Update solution vector x
  rnew <- r - c(alpha) * (A %*% p) # new residual vector
  znew <- solve(M, rnew) # update z
  beta <-
    crossprod(znew, rnew) / crossprod(z, r) # calculate beta coefficient
  p <- znew + c(beta) * p # update direction vector
  r <- rnew # update residual vector
  z <- znew # update z
  j <- j + 1 # update iteration counter

  conv[j] <- norm(r, "2") / norm(b, "2") # track convergence

  if (all(b == 0)) {
    # if b is a zero vector, track error
    err[j] <- norm(x, "I")
  }
  # Iteration counter:
  cat(sprintf('It. %3.0f: %20.16e\n', j, conv[j]), "\n")
}
cat("...finished!\n")
cat("The solution for x is:", x, "\n")

# Convergence plots (Convergence criterion over iterations):
plot(
  1:j,
  conv,
  type = "o",
  col = "blue",
  pch = 16,
  cex = 0.5,
  main = "Convergence Plot",
  xlab = "Iteration",
  ylab = "Convergence Criterion"
)
lines(1:j, conv, lty = 2, col = "blue")

plot(
  1:j,
  conv,
  type = "o",
  col = "red",
  pch = 16,

```

```

    cex = 0.5,
    main = "Convergence Plot",
    xlab = "Iteration",
    ylab = "Log Convergence Criterion"
)
lines(1:j, conv, lty = 2, col = "red")

# Return the convergence history and the final solution vector
return(list(conv = conv, x = c(x)))
}

```

Problem 2

Problem 2.1: Modified Lanczos that creates A-orthonormal bases

```

a_lanczos <-
function(A, # Input matrix
        v0, # Initial vector
        max_it = NULL, # Maximum number of iterations
        tol = 1e-10, # Tolerance value
        reorth = TRUE) { # Flag for reorthogonalization

  # Preliminary calculations and initializations: ----
  N <- length(v0)

  if(!is.null(max_it)) max_it <- N

  # Initialize the basis matrix V with zeros
  V <- matrix(0, nrow = N, ncol = max_it + 1)

  # Compute initial vector Av and its norm beta
  Av <- A(v0)
  beta <- sqrt(sum(v0 * Av))

  # Set the first basis vector
  V[, 2] <-
    v0 / beta # This is index 2 intentionally, 1 not returned

  # Normalize the initial vector Av
  Av <- Av / beta

  i <- 2 # Iteration counter

  # Lanczos iteration loop
  while (i < (max_it + 1) && beta > tol) {
    # Lanczos recursions
    w <- Av - beta * V[, i - 1]
    alpha <- sum(w * Av)
    w <- w - alpha * V[, i]

    # Reorthogonalization step

```

```

if (reorth) {
  # Subtract projections onto previous basis vectors (twice)
  w <- w - (V[, 2:i] %*% crossprod(V[, 2:i], A(w)))
  w <- w - (V[, 2:i] %*% crossprod(V[, 2:i], A(w)))
}

# Norm of New Vector Av
Av <- A(w)
beta <- sqrt(sum(w * Av))

# Store new vector if its norm is above the tolerance
if (beta > tol) {
  i <- i + 1
  Av <- Av / beta
  V[, i] <- w / beta
}
}

# Return the matrix V containing the A-orthonormal basis vectors
return(V[, 2:i])
}

```

Problem 2.2: Bayesian Conjugate Gradient Method

```

bayescg <- function(A, b, x, Sig, max_it = NULL, tol = 1e-6, delay = NULL,
  reorth = TRUE, NormA = NULL, xTrue = NULL,
  SqrtSigTranspose = NULL) {

  ## Variable definitions: ----

  # Size of the system
  N <- length(x)

  # Default Maximum Iterations
  if (is.null(max_it)) {
    max_it <- N
  }

  # Residual and first search direction
  r <- matrix(0, nrow = N, ncol = max_it + 1)
  r[, 1] <- b - A(x)
  S <- r

  # Inner products
  rIP <- numeric(max_it + 1)
  rIP[1] <- sum(r[, 1] * r[, 1])
  sIP <- numeric(max_it)

  # Array holding matrix-vector products
  SigAs_hist <- matrix(0, nrow = N, ncol = max_it)

```

```

# Convergence information
# If xTrue is supplied, more information is computed
rNorm <- sqrt(rIP[1])
Res2 <- numeric(max_it + 1)
if (is.null(NormA) || is.null(xTrue)) {
  bNorm <- norm(b, type = "2")
  Res <- rNorm / bNorm
  Res2[1] <- Res
}

if (!is.null(xTrue)) {
  xNorm <- norm(xTrue, type = "2")
  err_hist <- numeric(max_it + 1)
  err_hist[1] <- sum((x - xTrue) * A(x - xTrue))
  if (!is.null(NormA)) {
    xNormANorm <- norm(xTrue, type = "2") * NormA
    Res <- rNorm / xNormANorm
    Res2[1] <- Res
  }
  Res3 <- Res2
  tr_hist <- numeric(max_it + 1)
  tr_hist[1] <- sum(diag(A(Sig(diag(N)))))
}

i <- 0

## Iterating Through Bayesian Conjugate Gradient: ----

while (i < max_it && (is.null(tol) || Res > tol)) {
  # print(paste("Iteration:", i + 1))
  # Compute Matrix Vector Products
  As <- A(S[, i + 1])
  if (!is.null(SqrtSigTranspose)) {
    SigAs_hist[, i + 1] <- SqrtSigTranspose(As)
    SigAs <- Sig(SigAs_hist[, i + 1])
  } else {
    SigAs_hist[, i + 1] <- Sig(As)
    SigAs <- SigAs_hist[, i + 1]
  }
  ASigAs <- A(SigAs)

  # Search Direction Inner Product
  sIP[i + 1] <- abs(sum(S[, i + 1] * ASigAs))

  # Calculate next x
  alpha <- rIP[i + 1] / sIP[i + 1]
  x <- x + alpha * SigAs

  # Calculate New Residual
  r[, i + 2] <- r[, i + 1] - alpha * ASigAs

  if (reorth) {

```

```

# Reorthogonalize Residual
r_ip_inv <- 1 / rIP[1:(i + 1)]
ortho_term <- r[, 1:(i + 1)] %*% (t(r[, 1:(i + 1)]) %*% r[, i + 2])
diag_r_ip_inv <- diag(r_ip_inv, nrow = length(r_ip_inv), ncol = length(r_ip_inv))
if (ncol(ortho_term) == nrow(diag_r_ip_inv)) {
  r[, i + 2] <- r[, i + 2] - ortho_term %*% diag_r_ip_inv
} # else {
  # if dimension checking is not successful, give warning
  # warning("Dimensions of ortho_term and diag(r_ip_inv) do not match")
# }
}

# Compute Residual Norms
rIP[i + 2] <- sum(r[, i + 2] * r[, i + 2])
rNorm <- sqrt(rIP[i + 2])
if (!is.null(xTrue)) {
  err_hist[i + 2] <- sum((x - xTrue) * A(x - xTrue))
  tr_hist[i + 2] <- tr_hist[i + 1] - sum(diag(A(outer(SigAs, SigAs)))) / sIP[i + 1]

  rTrueNorm <- norm(b - A(x), type = "2")
  if (!is.null(NormA)) {
    Res <- rNorm / xNormANorm
    Res3[i + 2] <- rTrueNorm / xNormANorm
  } else {
    Res3[i + 2] <- rTrueNorm / bNorm
  }
}
if (is.null(NormA)) {
  Res <- rNorm / bNorm
} else if (is.null(xTrue)) {
  Res <- rNorm / NormA / norm(x, type = "2")
}
Res2[i + 2] <- Res

# Calculate next search direction
beta <- rIP[i + 2] / rIP[i + 1]
S[, i + 2] <- r[, i + 2] + beta * S[, i + 1]

i <- i + 1
}

## Return results: ----

info <- list(res = Res2[1:(i + 1)], search_dir = (sIP[1:i] ^ (-1/2)) * S[, 1:i])

if (!is.null(xTrue)) {
  info$actual_res <- Res3[1:(i + 1)]
  info$err <- err_hist[1:(i + 1)]
  info$trace <- tr_hist[1:(i + 1)]
}

if (!is.null(delay)) {
  delay <- min(delay, i)
}

```

```

    post_scale <- sum(rIP[(i - delay + 1):i]^2 / sIP[(i - delay + 1):i])
    info$scale <- post_scale
  }

  return(list(x = x, SigAs_hist = (sIP[1:i] ^ (-1/2)) * SigAs_hist[, 1:i], info = info))
}

```

Examples

```

# Set seed for reproducibility
set.seed(123)

# Number of rows in linear system of equations
n <- 10

# Create a random symmetric positive-definite coefficient matrix A
createPSDmatrixA <-
  function(n) {
    A <- matrix(rnorm(n * n), n, n)
    A <- crossprod(A) + n * diag(n)
  }
A <- createPSDmatrixA(n)

# Generate a random vector b
b <- rnorm(n)

# Initial guess for x
x0 <- rep(0, n)



# Find solution using Conjugate Gradient Algorithm (Problem 1.1) ----
cg_results <- cg(A, b, x0)

## Starting Conjugate Gradient algorithm...
## It.   1: 3.6631829446286168e-01
##
## It.   2: 1.0921528200647702e-01
##
## It.   3: 2.3227855416896950e-02
##
## It.   4: 5.1040155887402379e-03
##
## It.   5: 1.5745767128965691e-03
##
## It.   6: 3.2836615942801904e-04
##
## It.   7: 2.5295829214168488e-05
##
## It.   8: 4.2926878954089621e-06
##
## It.   9: 8.9040133058140012e-08
##

```



```
## It. 10: 2.6955219863623281e-16
##
## ...finished!
## The solution for x is: -0.03626387 0.0320711 -0.008359597 -0.02779354 -0.06587811 0.006584234 -0.0463
```

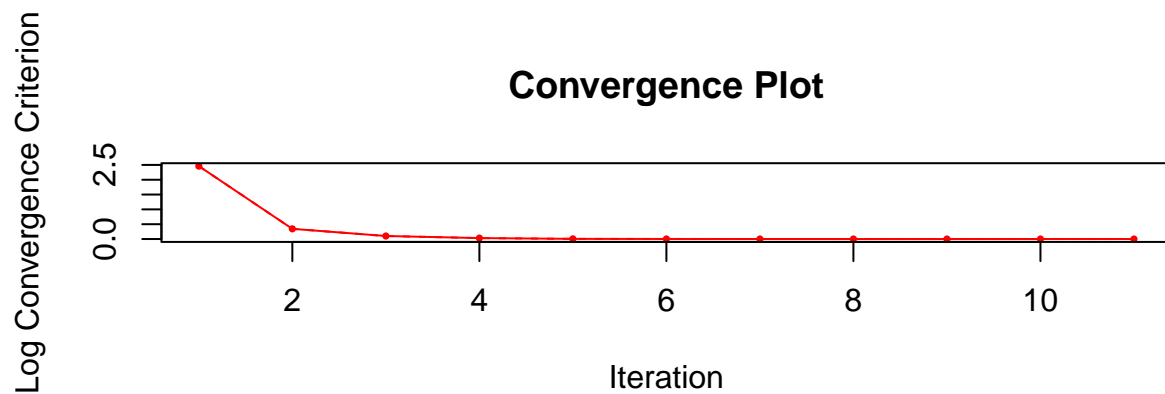



```
# Find solution using Preconditioned Conjugate Gradient Algorithm (Problem 1.2) ----
M1 <- diag(diag(A)) # Jacobi preconditioner
M2 <- diag(n) # Identity matrix

pcg_results1 <- pcg(A, b, x0, M1)
```

```
## Starting Preconditioned Conguate Gradient algorithm...
## It. 2: 3.4342927025576309e-01
##
## It. 3: 1.0103660141345808e-01
##
## It. 4: 3.1265737940853695e-02
##
## It. 5: 5.0008667290225781e-03
##
## It. 6: 1.4413296562075196e-03
##
## It. 7: 1.7965147376705902e-04
##
## It. 8: 5.2519637474155090e-05
```

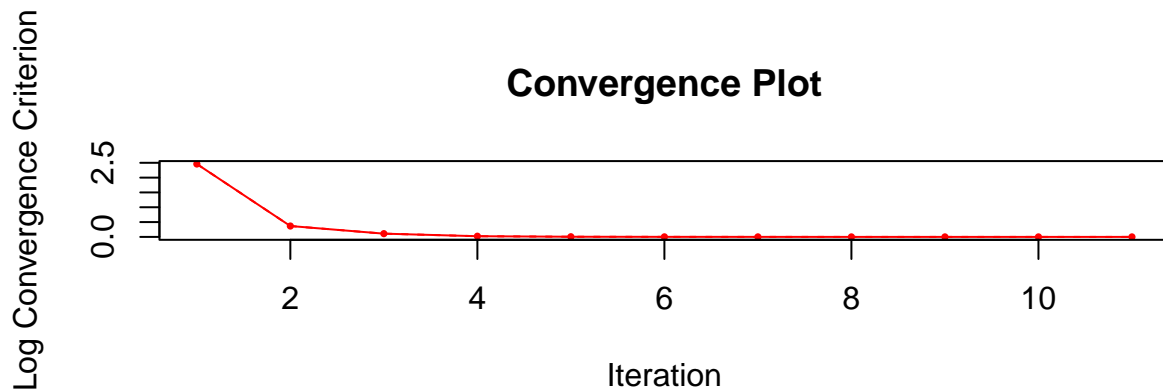

```
##
## It.   9: 1.8864624590058102e-06
##
## It.  10: 3.0535807772870021e-07
##
## It.  11: 3.2476221894550656e-17
##
## ...finished!
## The solution for x is: -0.03626387 0.0320711 -0.008359597 -0.02779354 -0.06587811 0.006584234 -0.0463
```



```
pcg_results2 <- pcg(A, b, x0, M2)
```

```
## Starting Preconditioned Conjugate Gradient algorithm...
## It.   2: 3.6631829446286168e-01
##
## It.   3: 1.0921528200647702e-01
##
## It.   4: 2.3227855416896950e-02
##
## It.   5: 5.1040155887402379e-03
##
## It.   6: 1.5745767128965691e-03
##
## It.   7: 3.2836615942801904e-04
##
```

```
## It.   8: 2.5295829214168488e-05
##
## It.   9: 4.2926878954089621e-06
##
## It.  10: 8.9040133058140012e-08
##
## It.  11: 2.6955219863623281e-16
##
## ...finished!
## The solution for x is: -0.03626387 0.0320711 -0.008359597 -0.02779354 -0.06587811 0.006584234 -0.0463
```



```
# Testing a_lanczos function (Problem 2.1) ----
Sig <- function(v) v # Define the preconditioner function Sig as the identity function
A_func <- function(v) A %*% v # Define the function A(v) to apply the matrix A to a vector v
v0 <- rnorm(n)
lanczos_vectors <- a_lanczos(A_func, v0, max_it = 10, tol = 1e-10, reorth = TRUE)
print("Lanczos vectors:")
```

```
## [1] "Lanczos vectors:"
```

```
lanczos_vectors
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] -0.051841140 0.101155612 0.06783553 -0.085010108 -0.0070868978
```

```
## [2,] 0.054780099 0.036835833 0.01801885 0.135973765 -0.0299798141
## [3,] -0.145777922 0.038424041 -0.06509586 0.024184744 -0.0009181462
## [4,] -0.005006363 0.027183998 0.05373108 -0.142822157 0.0696802424
## [5,] 0.046800737 0.023321466 -0.10005378 -0.067802672 -0.0511211261
## [6,] 0.027135163 0.123270156 -0.13412354 0.091139778 0.0648932751
## [7,] 0.009521862 -0.071981089 0.06618394 0.027178599 0.1347912637
## [8,] -0.057730261 0.110521848 0.02309881 0.052377238 -0.0571474350
## [9,] -0.076561875 0.064939400 -0.05725563 -0.007506557 0.1291959813
## [10,] -0.092278239 0.004469559 -0.04433263 0.091370324 -0.0846999523
##      [,6]      [,7]      [,8]      [,9]     [,10]
## [1,] 0.02322114 -0.201846517 0.03549281 0.025781726 0.009292058
## [2,] -0.01916601 -0.021334494 -0.17278404 -0.036957540 0.058224245
## [3,] 0.08010360 0.005285938 -0.07371684 0.002812149 0.157201242
## [4,] -0.11216775 0.125953413 -0.03405200 -0.088633611 0.130300575
## [5,] 0.05867997 -0.068890568 -0.03696493 -0.152851145 -0.034916536
## [6,] -0.02468506 -0.011782276 0.10677914 0.015989472 0.092718176
## [7,] 0.03452520 -0.077879904 0.06930027 -0.148910802 0.024526129
## [8,] 0.06436904 0.128326535 0.04796898 -0.085780914 -0.086232906
## [9,] -0.06077777 0.016042388 -0.09816357 0.028176329 -0.184069967
## [10,] -0.10795388 -0.054664030 0.06705214 -0.093956532 -0.007310916
```

```
# Find solution using the Bayesian Conjugate Gradient Algorithm (Problem 2.2) ----
bayescg_results <- bayescg(A_func, b, x0, Sig, max_it = 10, tol = 1e-6, reorth = TRUE)
```

```
# Find solution using implemented R solver ----
R_solver_results <- solve(A, b)
```

```
# Compare results ----
cat('Final solutions for x:\n')
```

```
## Final solutions for x:
```

```
cat('... using Conjugate Gradient Algorithm:\n')
```

```
## ... using Conjugate Gradient Algorithm:
```

```
cg_results$x
```

```
## [1] -0.036263867 0.032071103 -0.008359597 -0.027793538 -0.065878106
## [6] 0.006584234 -0.046571887 -0.082618303 -0.033472677 0.018976913
```

```
cat('... using Preconditioned Conjugate Gradient Algorithm (Jacobi):\n')
```

```
## ... using Preconditioned Conjugate Gradient Algorithm (Jacobi):
```

```
pcg_results1$x
```

```
## [1] -0.036263867 0.032071103 -0.008359597 -0.027793538 -0.065878106
## [6] 0.006584234 -0.046571887 -0.082618303 -0.033472677 0.018976913
```

```
cat('... using Preconditioned Conjugate Gradient Algorithm (Identity):\n')
```

```
## ... using Preconditioned Conjugate Gradient Algorithm (Identity):
```

```
pcg_results2$x
```

```
## [1] -0.036263867  0.032071103 -0.008359597 -0.027793538 -0.065878106  
## [6]  0.006584234 -0.046571887 -0.082618303 -0.033472677  0.018976913
```

```
cat('... using Bayesian Conjugate Gradient Algorithm:\n')
```

```
## ... using Bayesian Conjugate Gradient Algorithm:
```

```
bayescg_results$x
```

```
## [1] -0.036263867  0.032071103 -0.008359597 -0.027793538 -0.065878106  
## [6]  0.006584234 -0.046571887 -0.082618303 -0.033472677  0.018976913
```

```
cat('... using R Solver:\n')
```

```
## ... using R Solver:
```

```
R_solver_results
```

```
## [1] -0.036263867  0.032071103 -0.008359597 -0.027793538 -0.065878106  
## [6]  0.006584234 -0.046571887 -0.082618303 -0.033472677  0.018976913
```

```
all.equal(cg_results$x, R_solver_results, tolerance = 1e-5)
```

```
## [1] TRUE
```

```
all.equal(pcg_results1$x, R_solver_results, tolerance = 1e-5)
```

```
## [1] TRUE
```

```
all.equal(pcg_results2$x, R_solver_results, tolerance = 1e-5)
```

```
## [1] TRUE
```

```
all.equal(bayescg_results$x, R_solver_results, tolerance = 1e-5)
```

```
## [1] TRUE
```